

COMP27112

Introduction to Visual Computing

Image Processing Lecture 4.1

Terence Morley
School of Computer Science
The University of Manchester

Objectives for this Lecture

**To develop an understanding of
Point-processing techniques**

- ▶ Image histograms
- ▶ Grey-level mapping
- ▶ Thresholding
- ▶ Geometrical transformations

Image Histogram

Region of a greyscale image

152	158	166	165	163	171	175	169
159	160	160	164	173	172	169	173
153	160	155	158	172	169	164	175
151	156	157	154	161	165	167	174
150	150	156	157	157	164	171	171
154	155	157	158	160	161	166	165
154	153	150	154	161	159	158	161
152	156	152	153	159	160	155	161

What is its histogram?

Image Histogram

Histogram of the previous
image region

I	Freq		
150	3	163	1
151	1	164	3
152	3	165	3
153	3	166	2
154	4	167	1
155	3	168	0
156	3	169	3
157	4	170	0
158	4	171	3
159	3	172	2
160	5	173	2
161	5	174	1
162	0	175	2

Image Histogram

Pseudocode to generate histogram

```
int histogram[256] = {0};  
  
for(int i = 0; i < pixels in img; i++)  
{  
    histogram[ img[i] ]++;  
}
```

Image Histogram

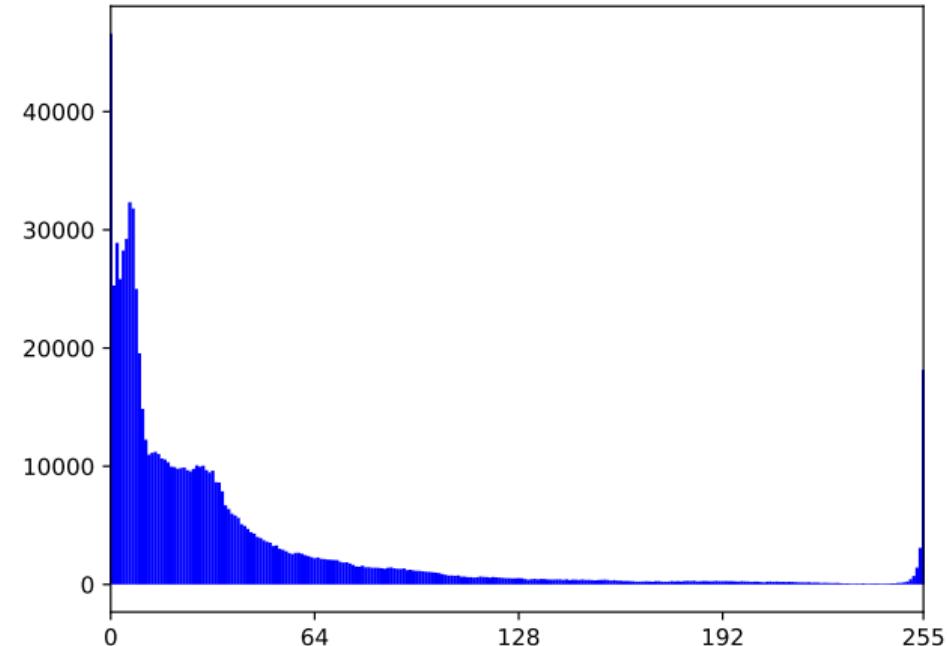


A *mostly* dark image

Image Histogram



Dark image

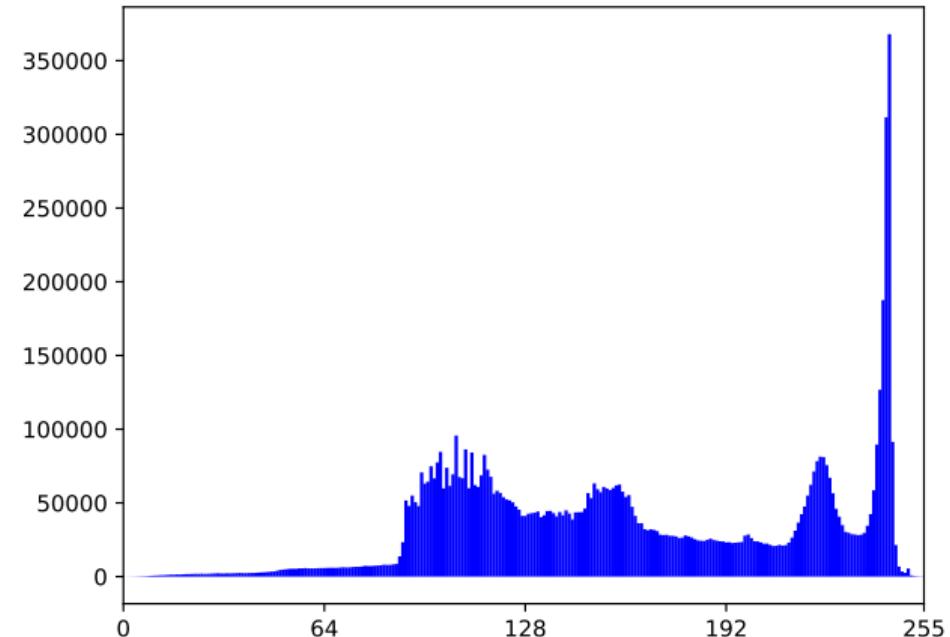


Histogram of dark image

Image Histogram



Light image



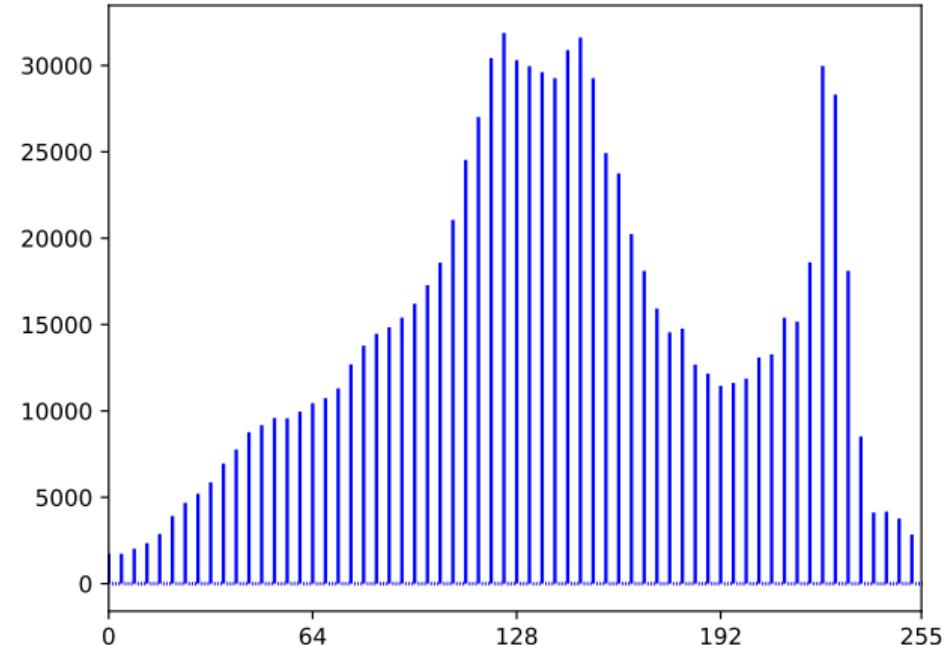
Histogram of light image

Image Histogram



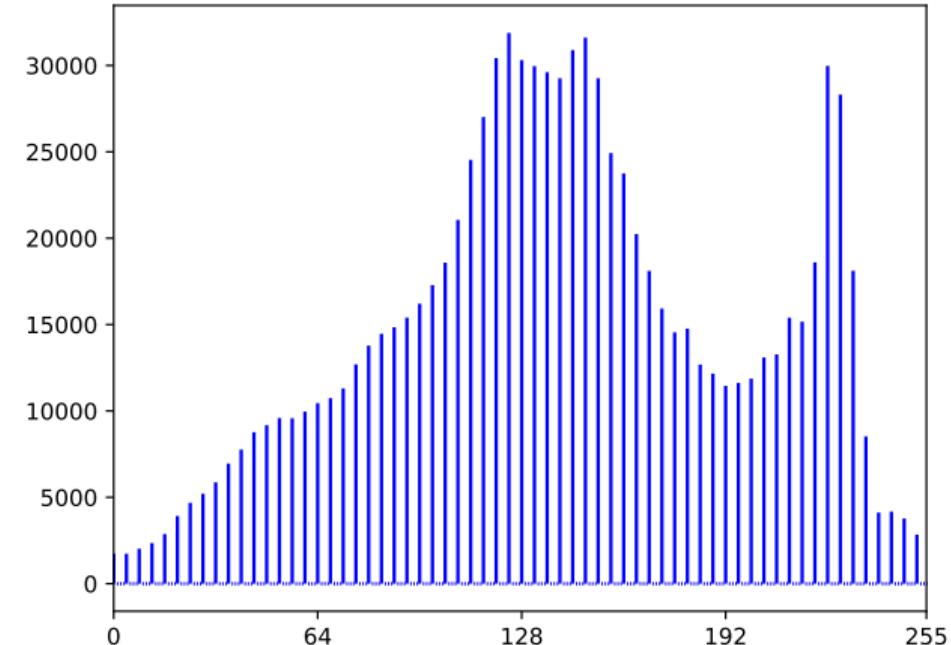
City scene

Image Histogram



Qu. Why does the histogram look like this?

Image Histogram



Histogram shows that only 64 grey levels exist in this image

Image Histogram

Incidentally, notice that the image with only 64 levels seemed perfectly fine.

Compare that with the same image having 32 levels (look at the sky):



Brightness Adjustment



$$I(x, y)$$

$$g$$

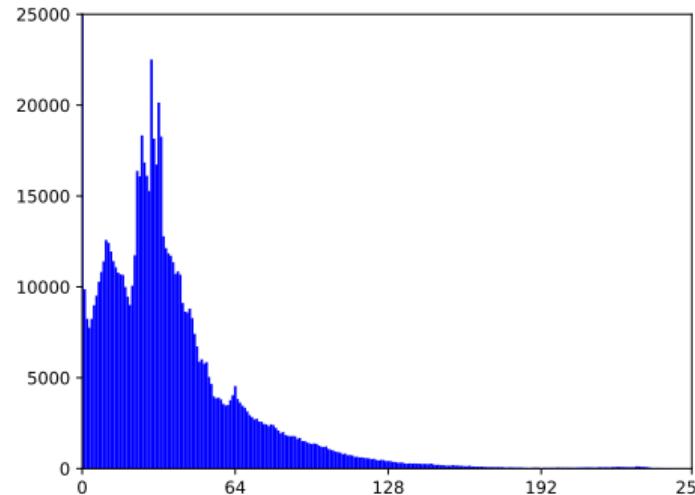


$$I'(x, y) = I(x, y) + k$$

$$g' = g + k$$

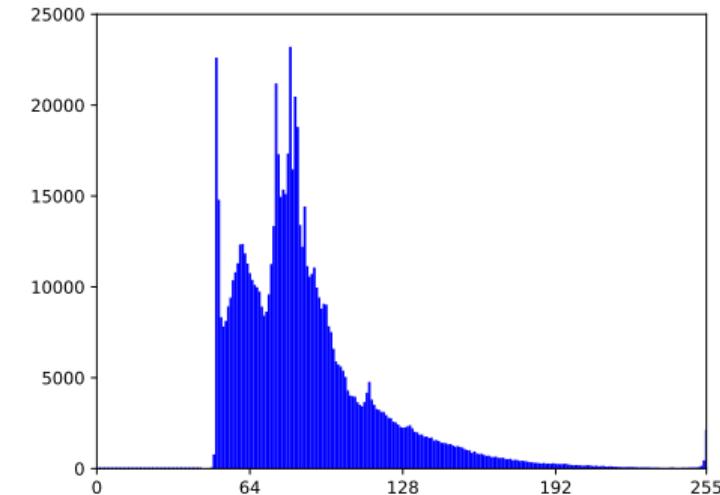
In this case $k = 50$

Brightness Adjustment



$$I(x, y)$$

$$g$$



$$I'(x, y) = I(x, y) + k$$

$$g' = g + k$$

In this case $k = 50$

Contrast Adjustment



$$I(x, y)$$

$$g$$

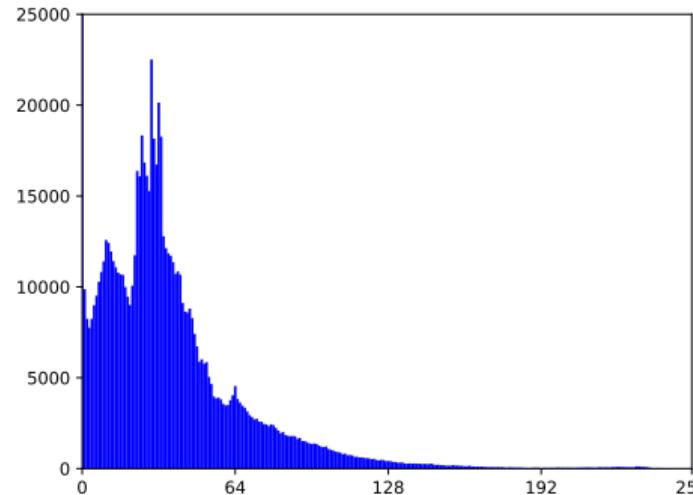
In this case $k = 1.8$



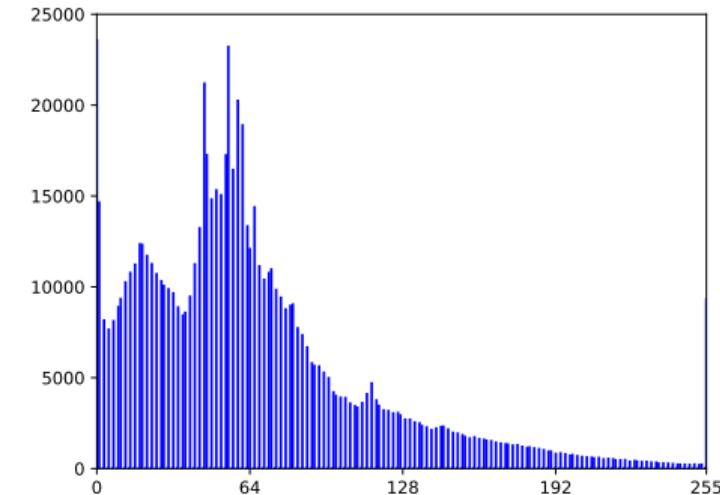
$$I'(x, y) = I(x, y) \times k$$

$$g' = g \times k$$

Contrast Adjustment



$$I(x, y)$$



$$I'(x, y) = I(x, y) \times k$$

$$g$$

$$g' = g \times k$$

In this case $k = 1.8$

Clipping

Qu. What is $255 + 1$ when dealing with unsigned bytes?

Clipping

Qu. What is $255 + 1$ when dealing with unsigned bytes?

$$255 + 1 = 0$$

$$255 + 2 = 1$$

etc.

Clipping

Qu. What is $255 + 1$ when dealing with unsigned bytes?

$$255 + 1 = 0$$

$$255 + 2 = 1$$

etc.

	1	1	1	1	1	1	1	1
+								1
1	0	0	0	0	0	0	0	0

(**1** = overflow)

Clipping

Qu. What is $255 + 1$ when dealing with unsigned bytes?

$$255 + 1 = 0$$

$$255 + 2 = 1$$

etc.

Qu. What effect will this have when adding 50 to an image to increase its brightness?

Clipping



Original



No clipping



Clipping

Clipping

When you process pixel values, you need to check for the result going over 255 and below 0

Example pseudocode

```
uint8 px = 250;  
...  
int32 p = px;  
p += 50; // Increase brightness  
if(p > 255) p = 255;  
if(p < 0) p = 0;  
px = p;
```

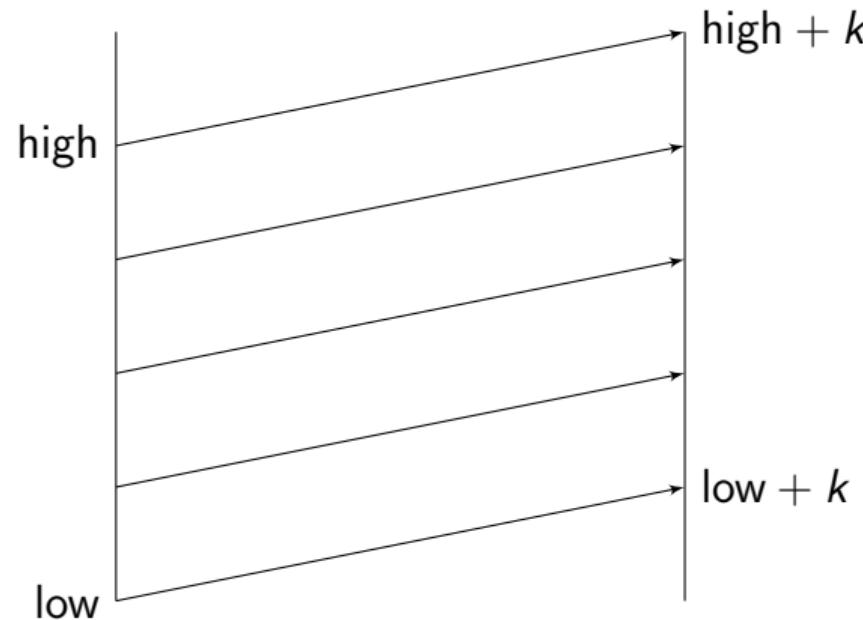
COMP27112

Introduction to Visual Computing

Image Processing Lecture 4.2

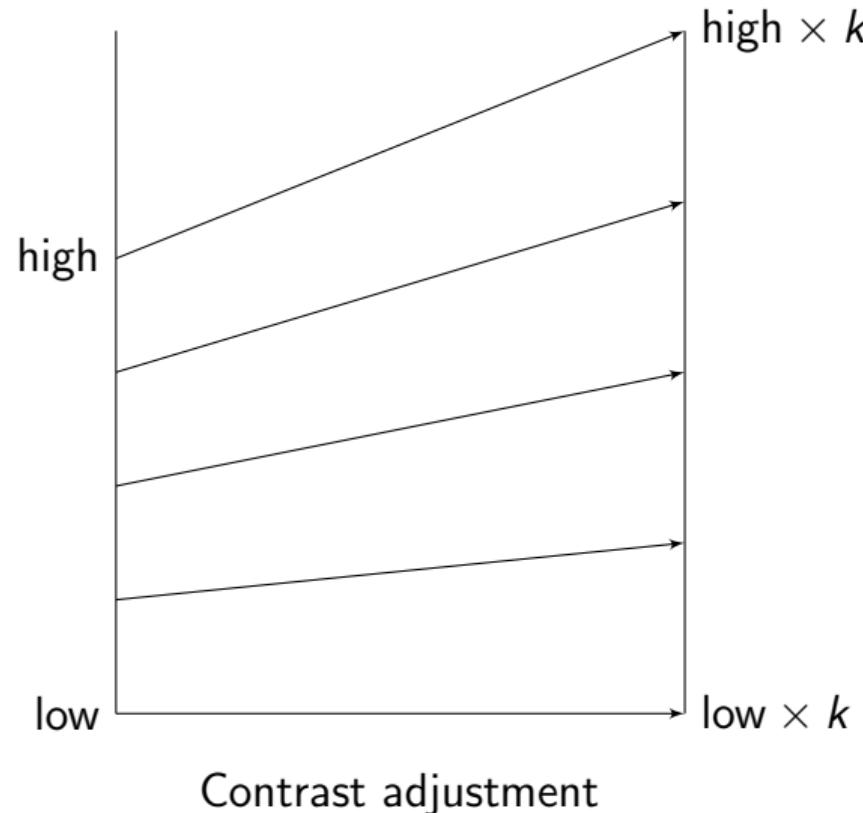
Terence Morley
School of Computer Science
The University of Manchester

Input-Output Mapping



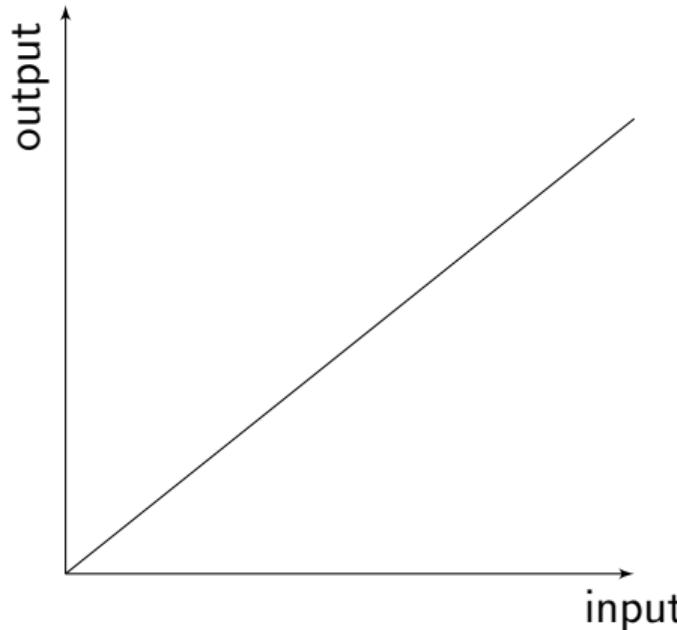
Brightness adjustment

Input-Output Mapping

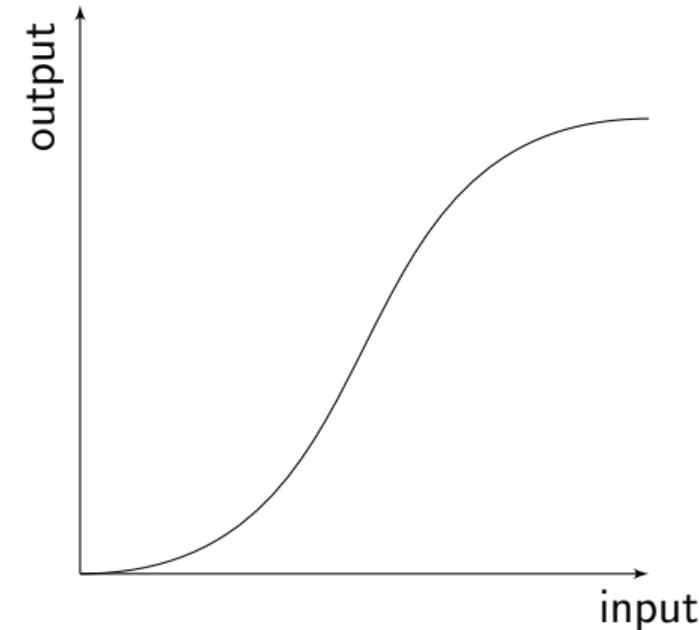


Input-Output Mapping

Creating lookup tables to relate an input grey level to an output grey level



Linear



Non-linear

Input-Output Mapping

Mappings can be created in code with a mathematical formula:

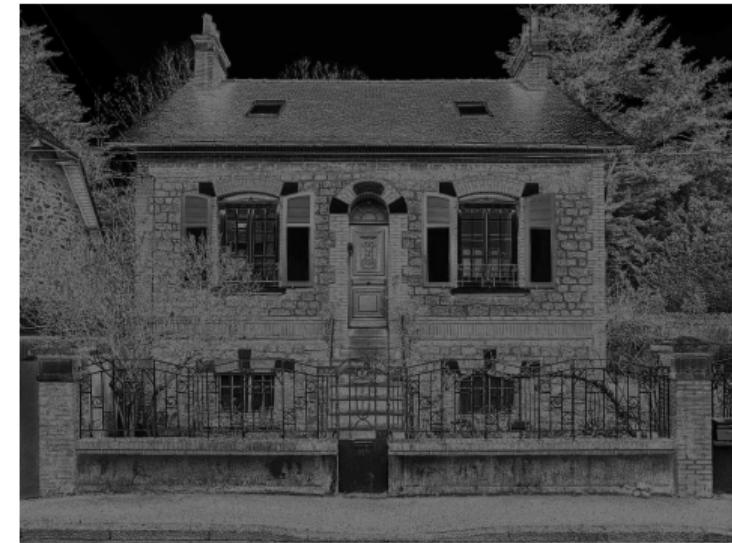
$$\text{img}[r][c] = \text{img}[r][c] * 1.5$$

Or with a lookup table:

map[256] = [0	2	3	5	6	8	9	11	12	14	...
--------------	---	---	---	---	---	---	---	----	----	----	-----

$$\text{img}[r][c] = \text{map}[\text{img}[r][c]]$$

Solarise

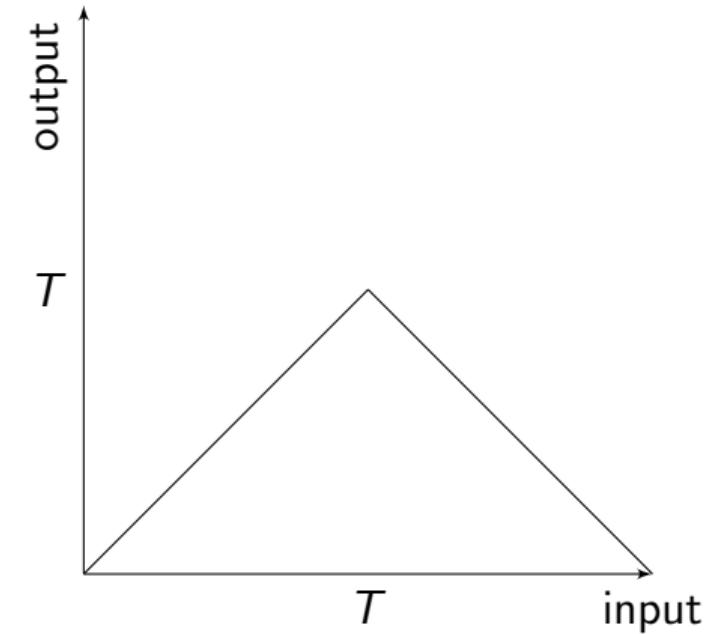


$$I(x, y)$$

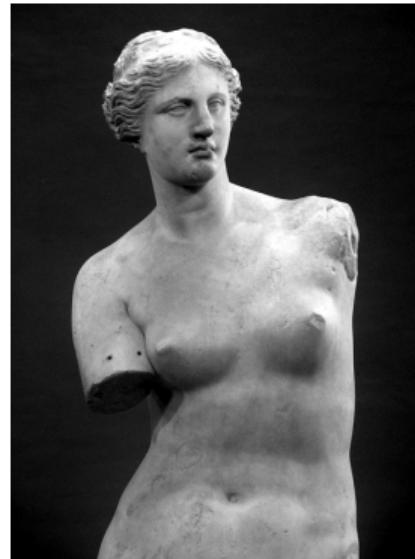
$$I'(x, y) = \begin{cases} 255 - I(x, y) & \text{if } I(x, y) > T \\ I(x, y) & \text{otherwise} \end{cases}$$

In this case $T = 127$

Solarise



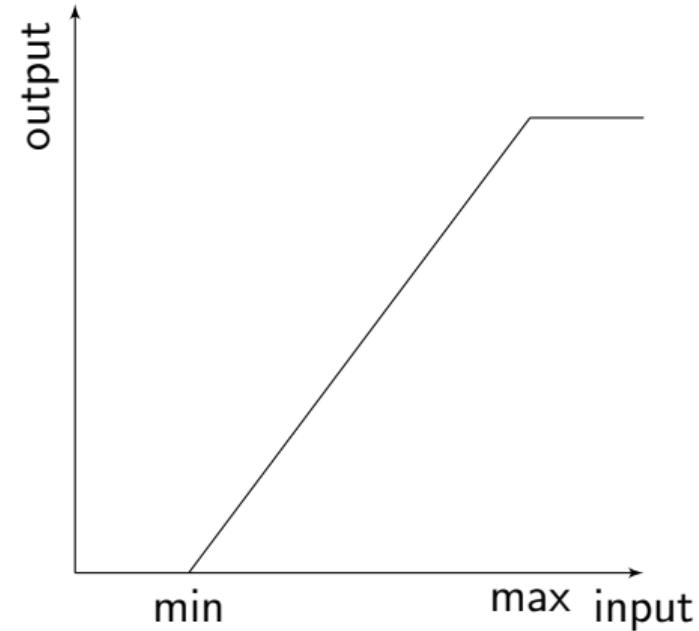
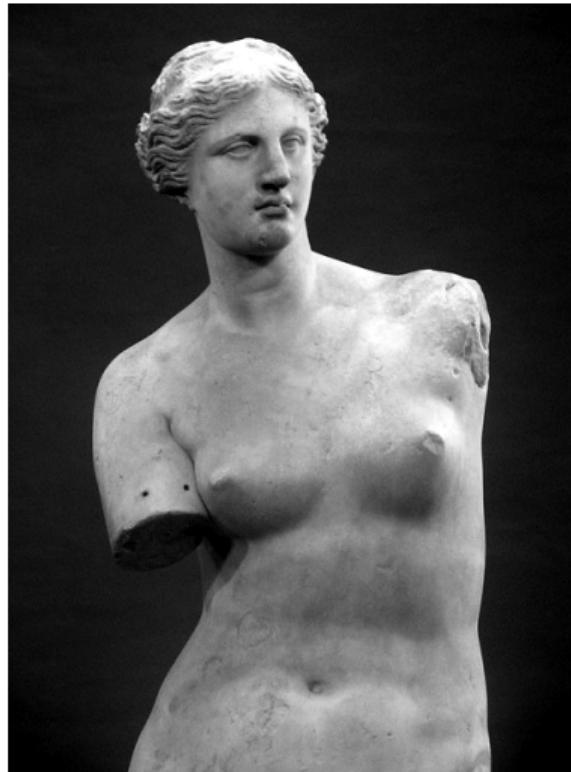
Min-Max Linear Stretch



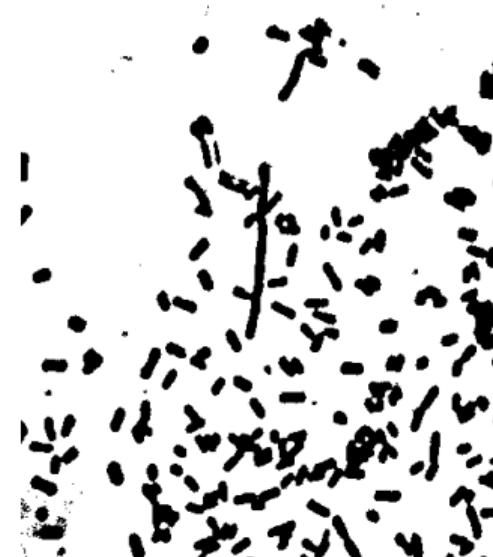
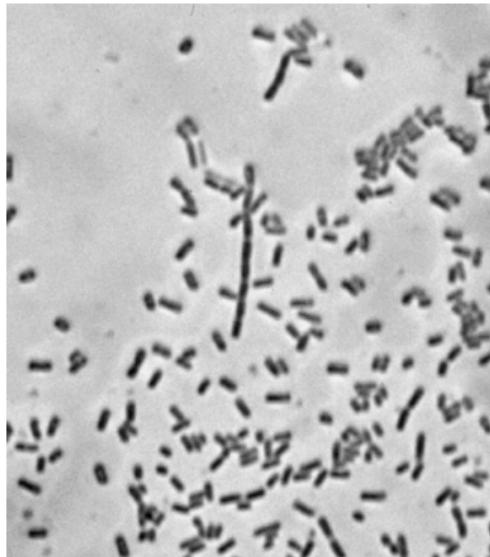
$$g' = \begin{cases} 0 & g < \min \\ \frac{g - \min}{\max - \min} 255 & \min \leq g \leq \max \\ 255 & g > \max \end{cases}$$

In this case, $\min = 50$ and $\max = 225$

Min-Max Linear Stretch



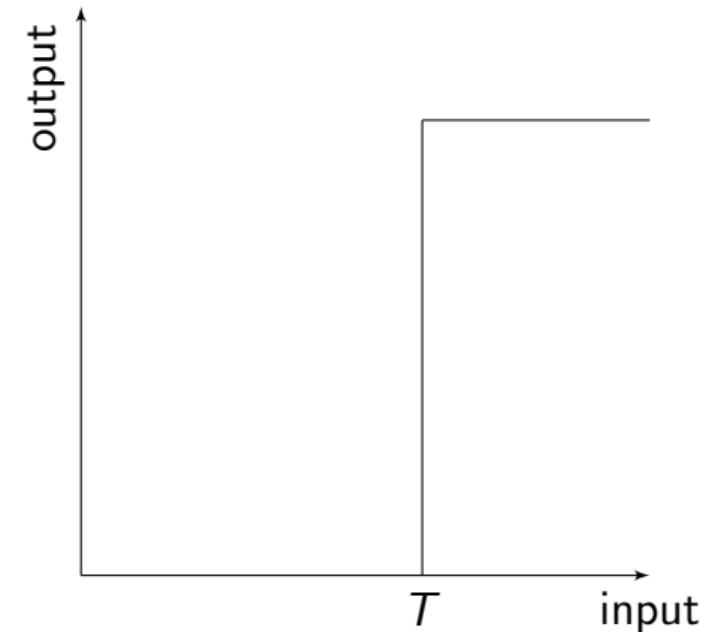
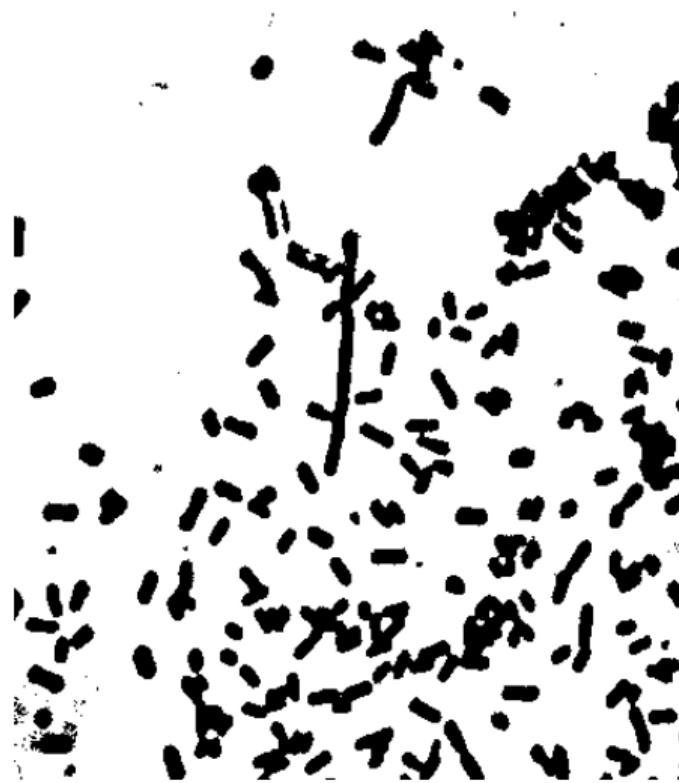
Thresholding



$$g' = \begin{cases} 255 & g \geq T \\ 0 & \text{otherwise} \end{cases}$$

Soil bacteria thresholded at $T = 170$

Thresholding

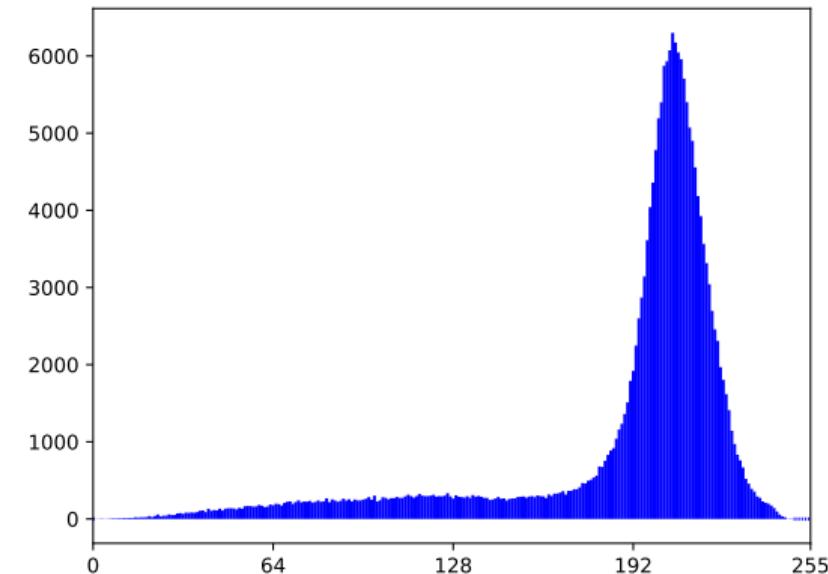
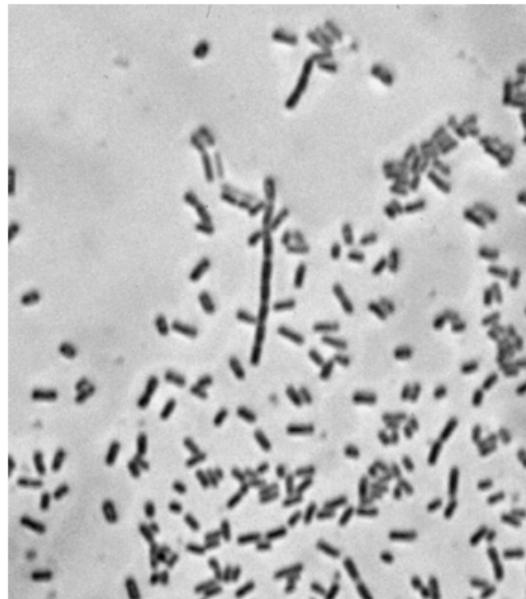


Thresholding

But how do we choose a value for the threshold, T ?

Thresholding

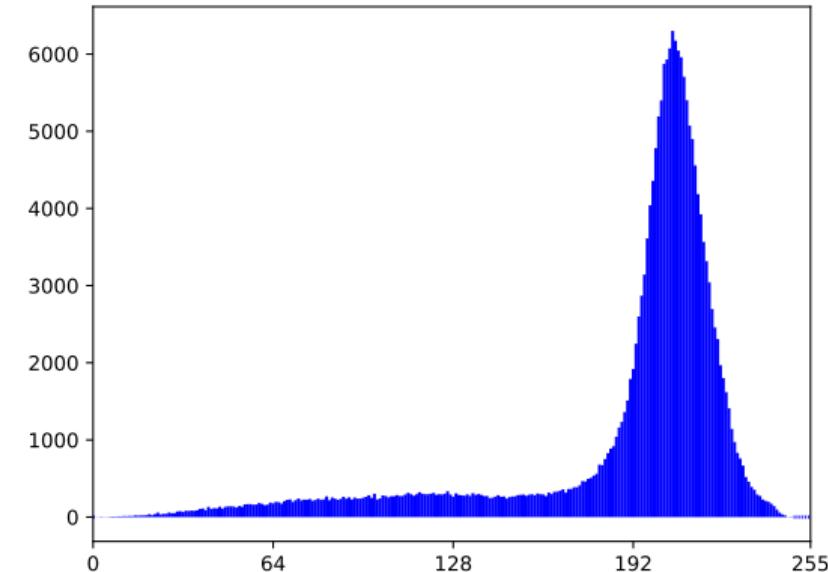
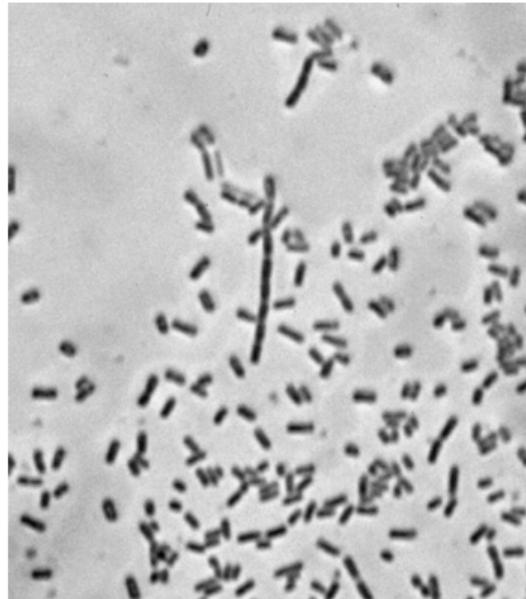
Manual selection of T



Qu. Where should we choose the threshold?

Thresholding

Manual selection of T



A threshold value of $T = 170$ was chosen after observation of the histogram

Automatic Thresholding

Manual selection of T requires a person to examine the histogram.

How can this be done automatically?

Automatic Thresholding

P-tile (percentile) calculation of T

If we know that an object should occupy a certain percentage of the pixels in an image, we can automatically select T

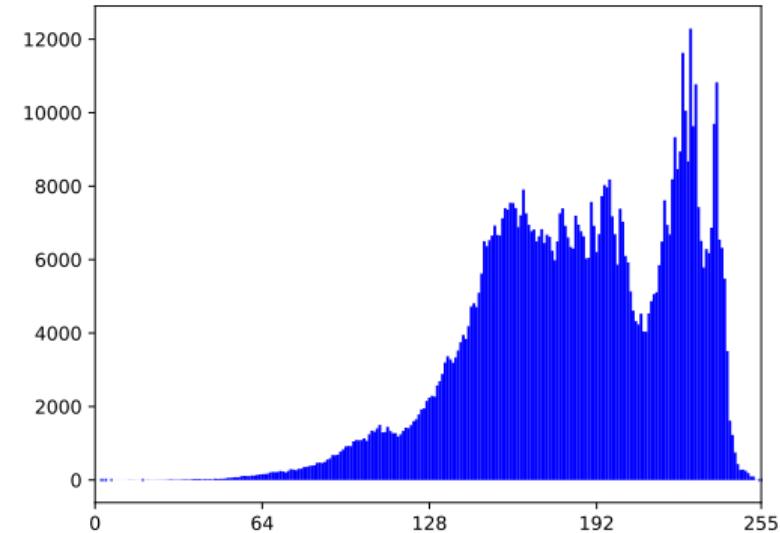
Example: a biscuit on a conveyor line should occupy 41% of the image.

Algorithm

- ▶ Calculate the number of pixels that should be object (% of image size)
- ▶ Create image histogram
- ▶ Accumulate frequencies until total exceeds number of expected object pixels
- ▶ Return current grey level as T

Automatic Thresholding

P-tile (percentile) calculation of T



Threshold calculated as $T = 176$ for object occupying 41% of the image

Automatic Thresholding

P-tile (percentile) calculation of T



Threshold calculated as $T = 176$ for object occupying 41% of the image

Automatic Thresholding

P-tile – Biscuit in low light

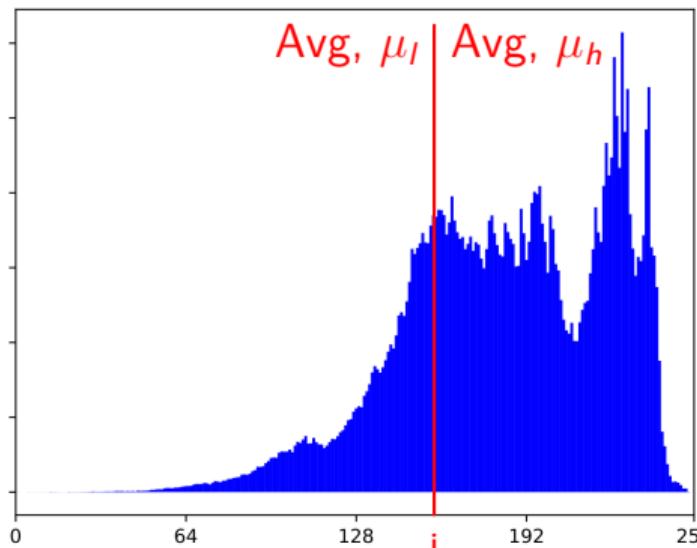


Threshold calculated as $T = 64$ for object occupying 41% of the image

Note Algorithm was not affected by the reduced lighting

Automatic Thresholding

Finding a threshold between the means of lower and upper parts of histogram



Algorithm

For $i \in [1, 255]$):

μ_l = Average intensity below i

μ_h = Average intensity above,
and including, i

$$\text{Midpoint}_i = (\mu_l + \mu_h)/2$$

Let $T = i$ where Midpoint_i
is closest to i

Automatic Thresholding

Simple method from previous slide

```
21 # Read image file (as greyscale, not colour)
22 filename = sys.argv[1]
23 img = cv.imread(filename, cv.IMREAD_GRAYSCALE)
24
25 # Find histogram - count of each intensity level 0 to 255
26 hist = cv.calcHist([img],[0],None,[256],[0,256])
27
28 min_diff = sys.float_info.max # a large value
29 T = 0
30 #
31 for i in range(1,256):
32     w0 = 0
33     w1 = 0
34     wt0 = 0
35     wt1 = 0
36     for a in range(0, i):
37         w0 += hist[a]
38         wt0 += (a * hist[a])
39     for b in range(i, 256):
40         w1 += hist[b]
41         wt1 += (b * hist[b])
42
43     if w0 == 0 or w1 == 0: continue
44
45     mw = 0.5 * ( (wt0/w0) + (wt1/w1))
46     diff = abs(mw - i)
47     if diff < min_diff:
48         min_diff = diff
49         T = i
50
51 print('Threshold =', T)
```

Python snippet

Automatic Thresholding

Otsu's method – similar to the previous method

Otsu's method [\[edit\]](#)

The algorithm exhaustively searches for the threshold that minimizes the intra-class variance, defined as a weighted sum of variances of the two classes:

$$\sigma_w^2(t) = \omega_0(t)\sigma_0^2(t) + \omega_1(t)\sigma_1^2(t)$$

Weights ω_0 and ω_1 are the probabilities of the two classes separated by a threshold t , and σ_0^2 and σ_1^2 are variances of these two classes.

The class probability $\omega_{0,1}(t)$ is computed from the L bins of the histogram:

$$\omega_0(t) = \sum_{i=0}^{t-1} p(i)$$

$$\omega_1(t) = \sum_{i=t}^{L-1} p(i)$$

For 2 classes, minimizing the intra-class variance is equivalent to maximizing inter-class variance:^[2]

$$\begin{aligned}\sigma_b^2(t) &= \sigma^2 - \sigma_w^2(t) = \omega_0(\mu_0 - \mu_T)^2 + \omega_1(\mu_1 - \mu_T)^2 \\ &= \omega_0(t)\omega_1(t)[\mu_0(t) - \mu_1(t)]^2\end{aligned}$$

which is expressed in terms of class probabilities ω and class means μ , where the class means $\mu_0(t)$, $\mu_1(t)$ and μ_T are:

$$\mu_0(t) = \frac{\sum_{i=0}^{t-1} ip(i)}{\omega_0(t)}$$

$$\mu_1(t) = \frac{\sum_{i=t}^{L-1} ip(i)}{\omega_1(t)}$$

$$\mu_T = \frac{\sum_{i=0}^{L-1} ip(i)}{\omega_0(t) + \omega_1(t)}$$

The following relations can be easily verified:

$$\begin{aligned}\omega_0\mu_0 + \omega_1\mu_1 &= \mu_T \\ \omega_0 + \omega_1 &= 1\end{aligned}$$

The class probabilities and class means can be computed iteratively. This idea yields an effective algorithm.

Algorithm [\[edit\]](#)

1. Compute histogram and probabilities of each intensity level
2. Set up initial $\omega_0(0)$ and $\mu_1(0)$
3. Step through all possible thresholds $t = 1, \dots$ maximum intensity
 1. Update ω_1 and μ_1
 2. Compute $\sigma_b^2(t)$
4. Desired threshold corresponds to the maximum $\sigma_b^2(t)$

Algorithm from Wikipedia – Looks complicated?

Automatic Thresholding

Otsu's method

```
25 # Find histogram - count of each intensity level 0 to 255
26 hist = cv.calcHist([img],[0],None,[256],[0,256])
27
28 hist /= (img.shape[0] * img.shape[1])
29
30 max_sigma_squared = 0
31 T = 0
32 #
33 for i in range(1,256):
34     sigma_squared = 0
35     w0 = 0
36     w1 = 0
37     wt0 = 0
38     wt1 = 0
39     for a in range(0, i):
40         w0 += hist[a]
41         wt0 += (a * hist[a])
42     for b in range(i, 256):
43         w1 += hist[b]
44         wt1 += (b * hist[b])
45
46     if w0 == 0 or w1 == 0: continue
47
48     sigma_squared = w0 * w1 * ( (wt0/w0) - (wt1/w1) )**2
49
50     if sigma_squared > max_sigma_squared:
51         max_sigma_squared = sigma_squared
52         T = i
53
54 print('Threshold =', T)
```

Python code for Otsu's method – Not complicated

Aside – Mathematical Notation

Complicated Mathematics

If you look at an image processing algorithm, it may sometimes look complicated

This is necessary to describe the method precisely

Spend some time working carefully through an algorithm
and it will usually turn out to be simpler than it seems

Automatic Thresholding

Otsu's method

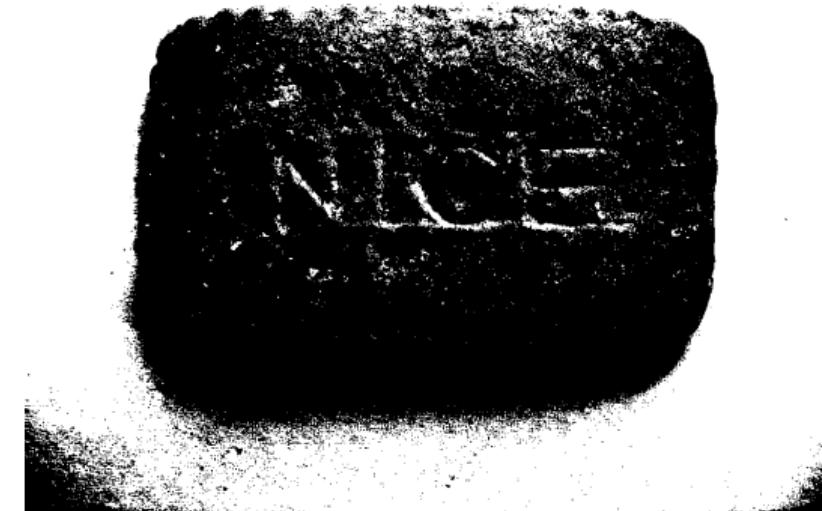
```
24 // Perform thresholding with a supplied threshold value
25 threshold(inputImage, outputImage, 128, 255, THRESH_BINARY);
26 // Save new image
27 imwrite("supplied_threshold.png", outputImage);
28
29 // Perform thresholding with Otsu threshold value
30 threshold(inputImage, outputImage, 0, 255, THRESH_BINARY | THRESH_OTSU);
31 // Save new image
32 imwrite("otsu_threshold.png", outputImage);
33
```

C code for thresholding with Otsu's method in OpenCV

Automatic Thresholding



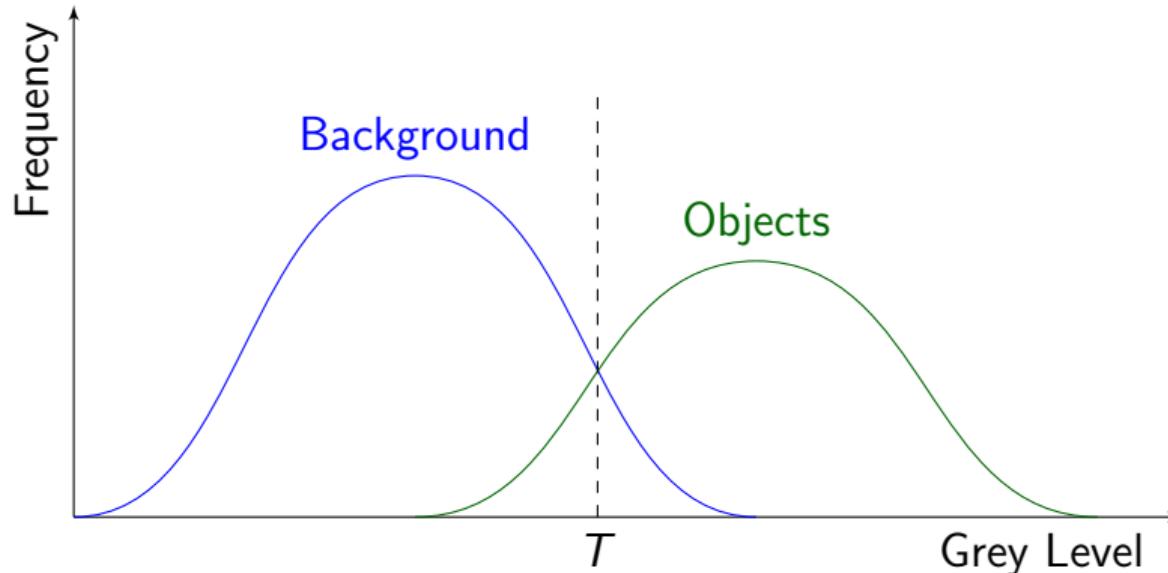
Using p-tile method (with 41%)



Using Otsu's method

Thresholding

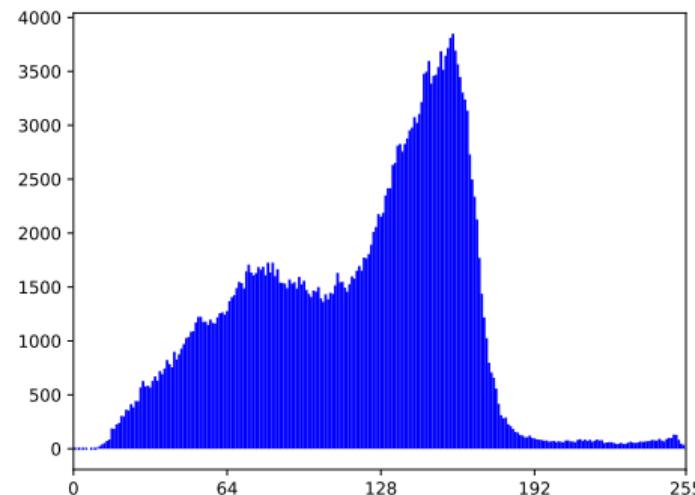
Misclassification



The histograms of background and objects overlap so some background pixels will be classified as objects and vice versa

Thresholding

Misclassification



COMP27112

Introduction to Visual Computing

Image Processing Lecture 4.3

Terence Morley
School of Computer Science
The University of Manchester

Geometrical Transformations



Spiralling succulent

Geometrical Transformations

Transformations Demo...



Geometrical Transformations

In the next few slides we will see affine transformations on this image from inside Manchester Museum



Scaling

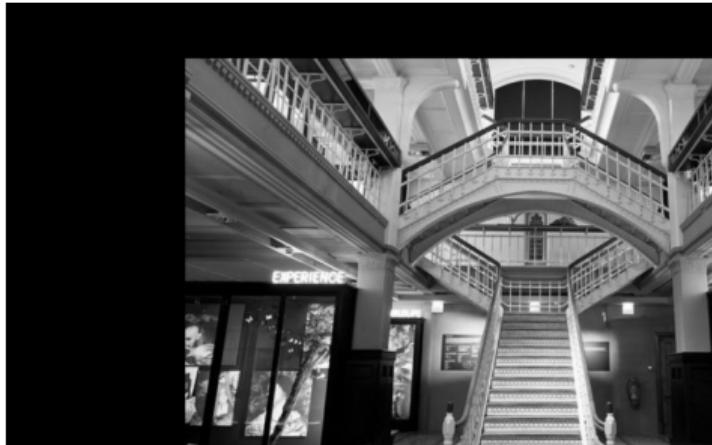


$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$
$$= \begin{bmatrix} s_x x \\ s_y y \end{bmatrix}$$

e.g. $s_x = 2, s_y = 2$
 $(1, 0) \rightarrow (2, 0)$
 $(2, 0) \rightarrow (4, 0)$

(We will discuss interpolation later)

Translation



$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$
$$= \begin{bmatrix} x + t_x \\ y + t_y \end{bmatrix}$$

e.g. $t_x = 160, t_y = 60$

$(1, 0) \rightarrow (161, 60)$

$(2, 0) \rightarrow (162, 60)$

Scaling and Translation

From the previous two slides we might perform a scaling followed by a translation as follows:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$
$$= \begin{bmatrix} s_x x + t_x \\ s_y y + t_y \end{bmatrix}$$

But this can be achieved with a single matrix multiplication...

Scaling and Translation

Using single matrix multiplication:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 & t_x \\ 0 & s_y & t_y \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} s_x x + t_x \\ s_y y + t_y \end{bmatrix}$$

Or, using **homogeneous** coordinates throughout which will allow us to multiply transformation matrices together:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & t_x \\ 0 & s_y & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} s_x x + t_x \\ s_y y + t_y \\ 1 \end{bmatrix}$$

Multiplication of Transformation Matrices

For example, we have a matrix for a translation and one for a scaling

We want to scale and then translate

We can multiply the matrices together (in reverse order) to obtain a single matrix to apply to the image

$$T = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & t_x \\ 0 & s_y & t_y \\ 0 & 0 & 1 \end{bmatrix}$$

Notice that the result is the same as the matrix on the previous slide

OpenCV Functions

Functions for performing geometrical transformations in OpenCV include:

warpAffine()
and
warpPerspective()

The former takes a 2×3 matrix and the latter a 3×3 matrix

Reflection/Flipping



$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$
$$= \begin{bmatrix} x \\ -y \\ 1 \end{bmatrix}$$

Note It is usual to translate the image centre to the origin, perform the transformation and then translate back. In this case for an image with height 480:

$$T = \begin{bmatrix} 1 & 0 & 0 \\ 0 & -1 & 480 \\ 0 & 0 & 1 \end{bmatrix}$$

Shear



$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & \lambda_x & 0 \\ \lambda_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$
$$= \begin{bmatrix} x + 0.8y \\ y \\ 1 \end{bmatrix}$$

In this case $\lambda_x = 0.8$, $\lambda_y = 0$

Remember It is usual to create a matrix to translate the image centre to the origin, perform a transformation and then translate back

Rotation



$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

For $\theta = 45^\circ$:

$$T = \begin{bmatrix} \sqrt{2}/2 & -\sqrt{2}/2 & 0 \\ \sqrt{2}/2 & \sqrt{2}/2 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Rotation

To rotate about **image centre**, create a matrix that translates centre to origin, rotates and translates back



$$R = \begin{bmatrix} \cos \theta & -\sin \theta & 0 \\ \sin \theta & \cos \theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Single matrix example for $\theta = 45^\circ$,
image = 640×480 :

$$T = \begin{bmatrix} \sqrt{2}/2 & -\sqrt{2}/2 & 263.5 \\ \sqrt{2}/2 & \sqrt{2}/2 & -156.0 \\ 0 & 0 & 1 \end{bmatrix}$$

Affine Transformations

Scale, Shear, Rotate, Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

As we have seen on previous slides, the different matrix entries are responsible for the following transformations:

- a, e → scaling
- b, d → shearing
- a, b, d, e → rotation
- c, f → translations

Camera Lens Distortion



Original



Barrel distortion

Camera Lens Distortion

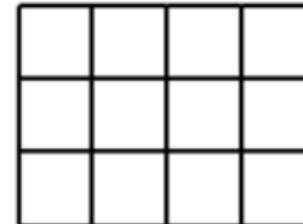


Original

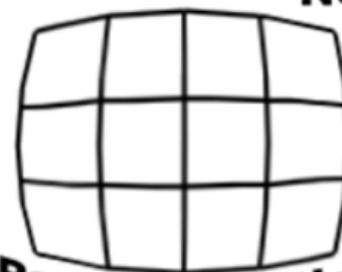


Pin-cushion distortion

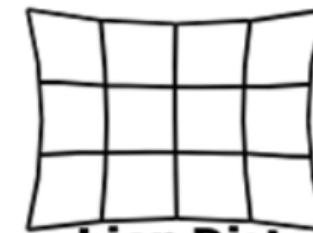
Camera Lens Distortion



No Distortion



Barrel Distortion



Pincushion Distortion

Camera Lens Distortion

Radial and Tangential distortions

Radial distortion

(r is the radius from the image centre, assuming here that the origin is the image centre)

$$x_{\text{corrected}} = x(1 + k_1 r^2 + k_2 r^4 + k_3 r^6 + \dots)$$
$$y_{\text{corrected}} = y(1 + k_1 r^2 + k_2 r^4 + k_3 r^6 + \dots)$$

Tangential distortion

$$x_{\text{corrected}} = x + (2p_1 xy) + p_2(r^2 + 2x^2)$$
$$y_{\text{corrected}} = y + p_1(r^2 + 2y^2) + (2p_2 xy)$$

Source: https://docs.opencv.org/3.4/d4/d94/tutorial_camera_calibration.html

Camera Lens Distortion

Radial and Tangential distortions

From the previous slide, OpenCV uses the constants in the **distortion matrix**:

$$[k_1 \quad k_2 \quad p_1 \quad p_2 \quad k_3]$$

Distortions can be corrected with the OpenCV function

`undistort()`

Camera Lens Distortion

Lens Distortion Demo...



Wide-angle Lens



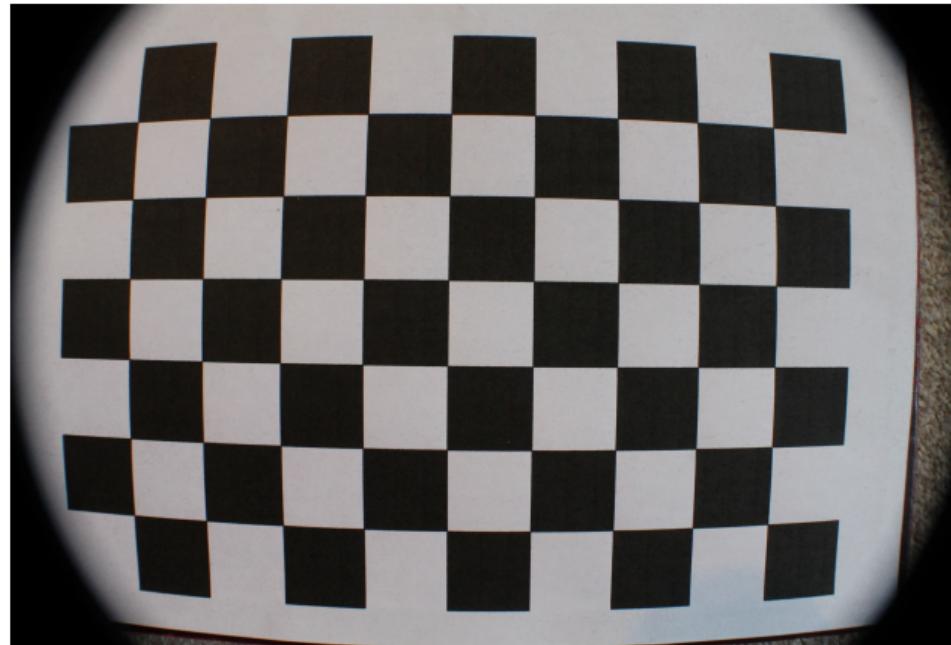
Petronas Towers, Malaysia. Image taken with a wide-angle lens

Fisheye Lens



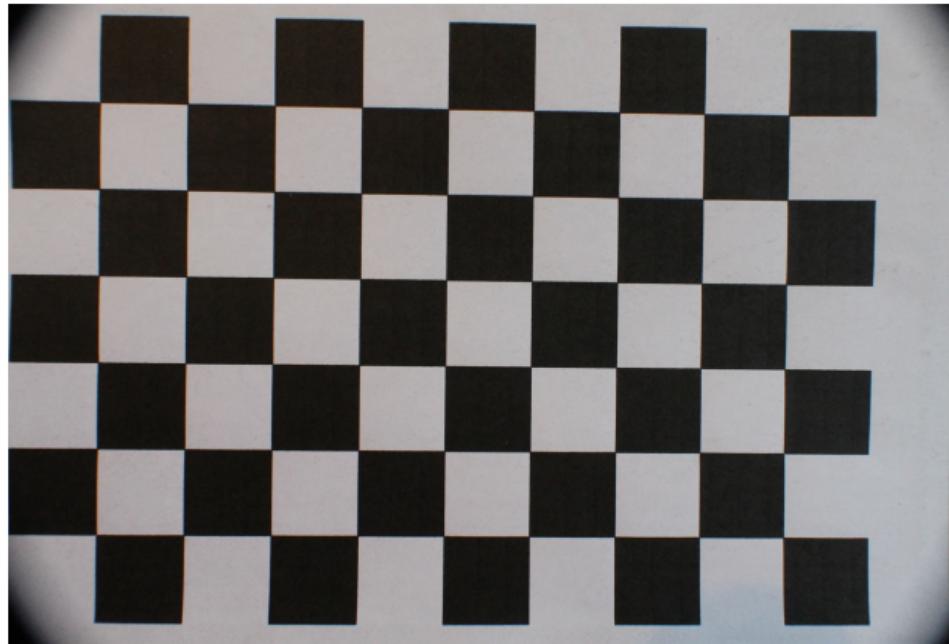
Barrel distortion from a fisheye lens
(This was taken with a cheap, low-quality, add-on fisheye lens)

Camera Calibration



The **chessboard pattern** that is used by OpenCV to allow you to correct for camera distortion

Camera Calibration



Chessboard image after calibration
(Used OpenCV to calculate a distortion matrix from multiple images)

Fisheye Lens



Barrel distortion from a fisheye lens

Fisheye Lens



Undistorted using matrix found by the chessboard calibration

COMP27112

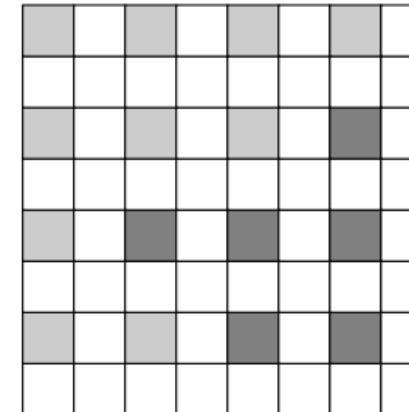
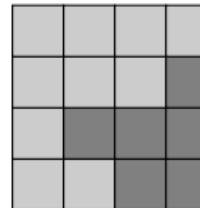
Introduction to Visual Computing

Image Processing Lecture 4.4

Terence Morley
School of Computer Science
The University of Manchester

Interpolation

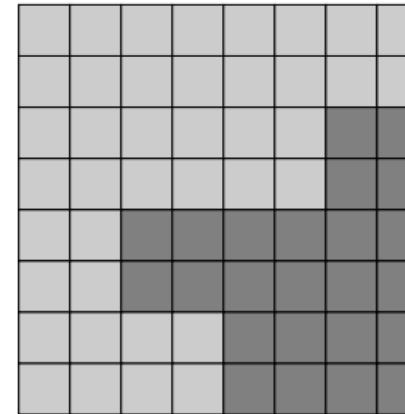
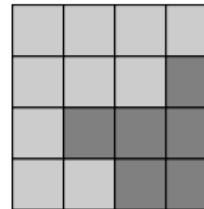
Scaling an image by a factor of 2



What about the missing pixels?

Interpolation

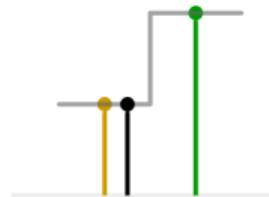
Scaling an image by a factor of 2 - Nearest neighbour interpolation



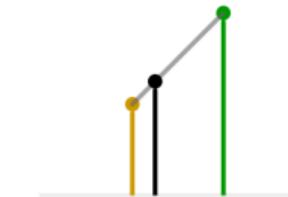
The result is a bit blocky/jagged.

Interpolation

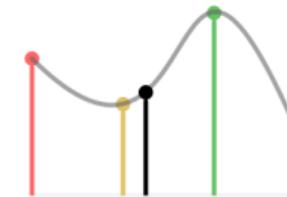
Some interpolation methods



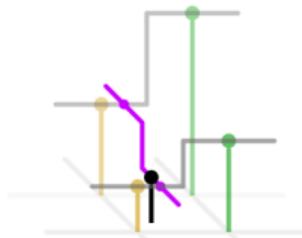
1D nearest-neighbour



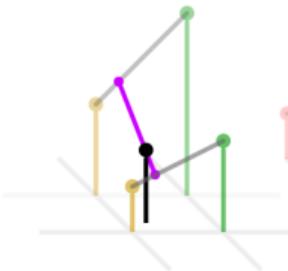
Linear



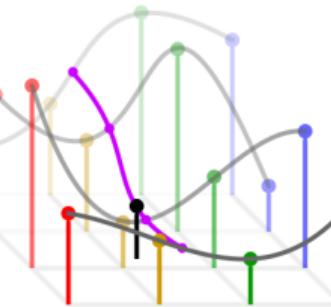
Cubic



2D nearest-neighbour



Bilinear



Bicubic

Forward/Reverse Mapping

In the previous **scaling** example, each pixel in the source image had somewhere to go in the destination image

For a **rotation**, source pixels might not ‘land’ in a destination pixel
(that is, coordinates can be non-integer)

If we round the coordinates, we will create holes in the destination image

Any ideas?

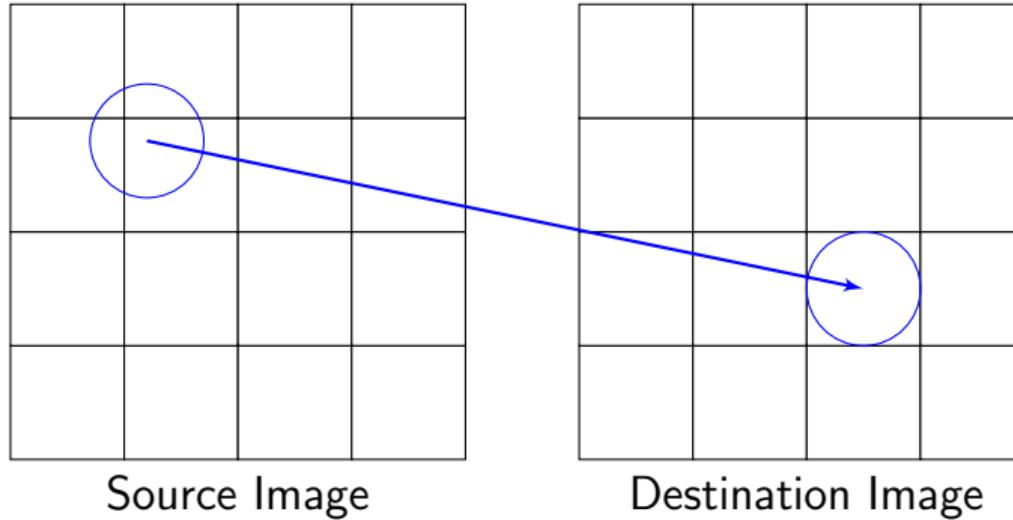
Forward/Reverse Mapping

Solution

Rather than cycling through source pixels and calculating a destination pixel,

cycle through destination pixels and calculate their values from source pixels...

Forward/Reverse Mapping



For each destination pixel, calculate where it came from
and then use interpolation to calculate the pixel value

Practical Problem

How might we (find and) level the horizon?



Summary

What have we discussed?

- ▶ Image histograms
- ▶ Grey-level mapping
- ▶ Thresholding
- ▶ Geometrical transformations